Benjamin Garcia

(626) 246-4778 | bentgarcia05@gmail.com | GitHub | LinkedIn | Portfolio

Education

University of California, Los Angeles — Los Angeles, CA
Bachelor of Science, Computer Science — *Expected Graduation: Dec. 2027*
Relevant Coursework: Data Structures and Algorithms, Operating Systems, Computer Architecture, Programming Languages

Technical Skills

Languages: TypeScript, Python, C++, Java, SQL, JavaScript
Frontend: React, Next.js, GraphQL, WebSockets, React Flow, Tailwind CSS
Backend & AI: Node.js, FastAPI, Flask, `asyncio`, Temporal, Redis, PostgreSQL, MySQL, MongoDB, OpenAI API
Infrastructure & Tools: Docker, AWS Lambda, Linux/Unix, Git, Playwright, Prisma, Zod

Experience

Production Engineer Fellow — Jun. 2026 – Present
Meta — *Remote*

- Built and deployed an open-source Flask/MySQL portfolio template to a **CentOS production server** using Python, Jinja, Docker, and Nginx, with DuckDNS routing and production environment configuration.
- Automated **CI/CD deployments** with `unittest`-based test gates, structured logging, and defensive error handling across deployment and runtime failures.
- Instrumented system and container health with **automated service checks, Prometheus metrics, Grafana dashboards, and alerting**.

Undergraduate Researcher — Dec. 2025 – Present
UCLA Human-Computer Interaction Research Lab — *Los Angeles, CA*

- Building a **human-in-the-loop AI research platform** that transforms scientific literature into abstract-grounded research Perspectives and controlled multi-agent deliberation, with researcher-gated updates to independently versioned documents.
- Evaluated two retrieval architectures across **5 domains and 20 uncached runs**, then implemented a hybrid combining position-aware, two-round Semantic Scholar retrieval with **SPECTER-v2 embeddings**, HDBSCAN clustering, DPP representative selection, and deterministic fallbacks.
- Engineered the **Next.js, FastAPI, and Supabase** platform with independently versioned documents and revision-checked persistence, backed by **173 pytest and 40 Playwright tests** spanning retrieval, interaction flows, and failure recovery.

Software Engineer Intern — Jun. 2026 – Aug. 2026
Lindy — *San Francisco, CA*

- Built Lindy's AI-agent guardrails and approvals platform with a policy store, arg-hashed approval tokens, fail-closed runtime enforcement, and Slack/web flows; retained **5,104 production decisions across 674 workspaces**.
- Took over a stalled Integrations unification effort and drove unified personal and team integrations to **100% GA**; shipped an eight-PR, feature-flagged foundation and a unification change that net-deleted **4,521 lines**.
- Led frontend and integration engineering for Lindy's existing-user migration to Teammate, shipping Slack authorization, durable memory import, a five-step wizard, and E2E launch gates; go-live reached **78 production completions and 50 activations**.
- Traced a silent 21-minute Slack hydration stall to KEDA scale-down and five code defects, then migrated the workflow from BullMQ to Temporal in one day; the following 14-day window recorded **409 runs with one failure**.

Software Engineer Intern — Mar. 2025 – Sep. 2025
Todd Agriscience — *Los Angeles, CA*

- Shipped production Next.js crop-intelligence dashboards for **5–10 clients**, integrating Palantir Foundry pipelines and AWS Lambda services to surface real-time agricultural insights.
- Built data pipelines visualizing **20+ agricultural streams** across soil, irrigation, weather, and yield data, enabling faster access to operational farm-management information.

Projects

PolicyC | *TypeScript, Python, asyncio, SQLite, OpenAI API, Zod, FastAPI* — Jul. 2026 – Present

- Built AI infrastructure that compiles request-specific policy context from a **43-node, six-domain dependency graph**, reducing mean model input by up to **98.23%** while preserving required policy dependencies.
- Engineered a resumable Python `asyncio` runtime for concurrent evaluation, reducing elapsed time by **72.2% across three held-out experiments** and achieving **3.60× effective throughput**.
- Ran **960 GPT-5 mini calls for $2.66** across paired held-out evaluations, preserving up to **86.49% of critical obligations** while projecting a **60.6% reduction** in uncached model-input cost versus full prompts.