Benjamin Garcia

(626) 246-4778 | bentgarcia05@gmail.com | GitHub | LinkedIn | Portfolio

Education

University of California, Los Angeles — Los Angeles, CA
Bachelor of Science, Computer Science — *Expected Graduation: Dec. 2027*
Relevant Coursework: Data Structures and Algorithms, Operating Systems, Computer Architecture, Programming Languages

Technical Skills

Languages: TypeScript, Python, C++, Java, SQL, JavaScript
Frontend: React, Next.js, GraphQL, WebSockets, React Flow, Tailwind CSS
Backend & AI: Node.js, FastAPI, Flask, `asyncio`, Temporal, Redis, PostgreSQL, MySQL, MongoDB, OpenAI API
Infrastructure & Tools: Docker, AWS Lambda, Linux/Unix, Git, Playwright, Prisma, Zod

Experience

Software Engineer Intern — Jun. 2026 – Present
Lindy — *San Francisco, CA*

- Shipped multi-account Email and Meetings infrastructure serving **27,252 users**, isolating OAuth credentials, settings, and agent state across accounts to prevent removal, reconnection, sync, and opt-out flows from silently re-enabling disabled automation.
- Owned account-management and authorization guardrails across web, mobile, and Lindy Teammate for a platform supporting **5,000+ integrations**, enforcing workspace permissions and secure tool access across add, reconnect, delete, rename, and sharing workflows.
- Improved production Gmail and Outlook agent experiences by integrating user instructions, learned writing style, holiday exclusions, account-aware classification, and rendering fixes, helping reduce scheduling-draft edits from **14.9% to 3.2%** and draft-alert pushback from **6.7% to 3.9%**.

Production Engineer Fellow — Jun. 2026 – Present
Meta — *Remote*

- Deployed and operated a containerized Flask and MySQL service on a remote CentOS server, configuring Docker workflows, DuckDNS routing, and production environment settings.
- Improved production reliability by automating **service health checks** and applying structured logging, defensive error handling, and systematic debugging to deployment and runtime failures.
- Designed resilient recovery strategies using **exponential backoff, circuit breakers, asynchronous queues, and CI/CD automation** to reduce manual intervention during service failures.

Undergraduate Researcher — Dec. 2025 – Present
UCLA Human-Computer Interaction Research Lab — *Los Angeles, CA*

- Building a **human-in-the-loop** AI research platform that transforms retrieved scientific literature into evidence-grounded agent perspectives, structured debate, and refined research hypotheses.
- Built a React Flow developer interface backed by FastAPI event streams, visualizing **27 workflow event types** across agent reasoning, critique, refinement, and persistence for inspectable AI workflows.
- Developed a **5-stage, 20-step** literature-grounded agent workflow that retrieves Semantic Scholar papers, clusters evidence by embeddings, synthesizes citation-grounded researcher personas, and surfaces provenance in the UI.

Software Engineer Intern — Mar. 2025 – Sep. 2025
Todd Agriscience — *Los Angeles, CA*

- Shipped production Next.js crop-intelligence dashboards for **5–10 clients**, integrating Palantir Foundry pipelines and AWS Lambda services to deliver real-time agricultural insights.
- Built data pipelines visualizing **20+ agricultural streams** across soil, irrigation, weather, and yield data, enabling faster access to operational farm-management information.

Projects

PolicyC | *TypeScript, Python, asyncio, SQLite, OpenAI API, Zod, FastAPI* — Jul. 2026 – Present

- Built AI infrastructure for compiling request-specific policy context from a 43-node, six-domain dependency graph, reducing mean model input by up to **98.23%** while preserving required policy dependencies.
- Engineered a resumable Python `asyncio` execution runtime that completed three held-out experiments in 37.24 aggregate minutes versus 134.11 minutes of cumulative provider-call latency, achieving **72.2% lower elapsed time** and **3.60× effective throughput**.
- Automated **960 GPT-5 mini calls** across paired held-out evaluations for **$2.66**, achieving up to **86.49% critical-obligation preservation** while projecting a **60.6% reduction** in uncached model-input cost versus full prompts.